




Executive Summary

The manuscript presents a novel approach to automatic question generation aligned with Bloom's Taxonomy using an integrated framework of Retrieval-Augmented Generation (RAG), Large Language Models (LLMs), and ModernBERT-Bloom-Taxonomy classifier. While the foundational elements are strong, there is room for improvement in explicit justification, detailed analysis, and narrative references to visual components.

Key Synthesis of Dialogue

Critic's Concerns:

- Lack of explicit linkage between objectives and specific challenges.
- Absence of real-world implications or case studies.
- Inadequate justification for technology choices (e.g., ModernBERT).
- Insufficient detailed analysis across different Bloom's levels.
- Missing narrative references to visual elements.

Advocate's Rebuttals:

- The manuscript implicitly supports alignment through detailed descriptions and rationales.
- Real-world implications are substantiated by the identified research gap.

- Justifications for technology choices are grounded in performance metrics and literature citations.
- Detailed analyses support claims about model performance.
- Visual components are well-documented but need explicit narrative references.

Point-by-point Audit

S1: Foundational Audit Analysis

Alignment Status: The objectives are well-aligned with the research problem, though explicit linkage to specific challenges is missing. **Suggestion for Improvement:** Explicitly link each objective to a specific challenge identified in the problem statement.

Research Gap: Clearly defined but lacks real-world implications or case studies. **Suggestion for Improvement:** Provide concrete examples and scenarios demonstrating the necessity of an integrated approach.

S2: METHODOLOGY

Methodological Synergy: The preprocessing pipeline is well-documented, but transitions between stages need more clarity. **Suggestion for Improvement:** Explain how cleaned data transitions into structured chunks for embedding.

Justification: Choice of ModernBERT lacks detailed benchmarks or comparisons. **Suggestion for Improvement:** Include specific performance metrics and comparisons with alternative models.

S3: Scientific Audit

Numerical & Mathematical Audit: Equations are consistent but lack detailed analysis across Bloom's levels. **Suggestion for Improvement:** Provide a breakdown of model performance at different cognitive levels.

Claim-to-Evidence Ratio: Claims about fine-tuned model performance need more detailed explanations. **Suggestion for Improvement:** Include specific analyses explaining variability in performance.

S4: Scientific Maturity

Interpretive Consequences: Logical links are strong but could benefit from deeper interpretation. **Suggestion for Improvement:** Provide a comprehensive analysis of how different cognitive levels impact model performance.

S5: Visual Audit

Narrative References: Tables and figures lack explicit narrative references in the body text. **Recommendation for Improvement:** Ensure each visual element has an explicit reference in the body text.

Final Verdict

Verdict: Accept with Minor Revisions

Key Revisions Required:

1. Explicitly link objectives to specific challenges identified.
2. Provide real-world implications and case studies demonstrating the necessity of the integrated approach.
3. Include detailed benchmarks or comparisons for technology choices (e.g., ModernBERT).
4. Break down model performance across different Bloom's levels with detailed analyses.
5. Ensure each visual element has an explicit narrative reference in the body text.

Top 3 Fatal Issues:

- Lack of explicit linkage between objectives and specific challenges.
- Absence of real-world implications or case studies.
- Inadequate justification for technology choices (e.g., ModernBERT).

The manuscript demonstrates strong foundational elements, innovative methodology, and robust data evidence. However, addressing the identified issues will significantly enhance its clarity and robustness, making it suitable for Q1 journal standards.

Actionability:

- Revise to explicitly link objectives with specific challenges.
- Include real-world implications and case studies.
- Provide detailed benchmarks or comparisons for technology choices.
- Break down model performance across different Bloom's levels.
- Ensure narrative references to all visual elements.

Foundational Audit Analysis

1. Research Problem & Objectives Alignment

Alignment Status: Strong alignment. The research problem is clearly defined and focuses on the limitations of existing automatic question generation systems in terms of cognitive level alignment with Bloom's Taxonomy. The objectives are well-aligned with addressing these limitations through an integrated framework that includes RAG, LLMs, and a pre-trained classifier.

Suggestion for Improvement: Ensure that each objective is directly linked to solving specific challenges identified in the problem statement. For example, explicitly state how the regenerative feedback loop addresses the issue of inconsistent cognitive level alignment.

2. Research Gap (Explicit vs Implicit)

Gap Identification: Explicitly defined. The manuscript clearly identifies a gap where existing systems lack an integrated approach for generating questions that are aligned with Bloom's Taxonomy levels and do not provide mechanisms to verify and refine these alignments post-generation.

Suggestion for Improvement: Provide more context on why this specific research gap is critical in the field of educational technology. Highlighting real-world implications or case studies where such a system could make significant improvements would strengthen the argument.

3. Rationale for Chosen Approach

Rationale: Strong rationale. The manuscript provides a clear rationale for integrating RAG, LLMs, and ModernBERT-Bloom-Taxonomy classifier to address the identified research gap. The approach is justified by its ability to generate contextually relevant questions with verified cognitive levels.

Suggestion for Improvement: Elaborate on why these specific technologies (RAG, LLMs, ModernBERT) were chosen over others. Discuss any potential alternatives and explain why they were not selected, further justifying the choice of this integrated framework.

4. Innovation Level

Innovation: Substantial innovation. The manuscript introduces a novel approach by integrating RAG with multiple LLMs and a pre-trained classifier to ensure cognitive level alignment in question generation. This is a significant step beyond existing incremental improvements.

Suggestion for Improvement: Highlight the unique contributions of this work more prominently, such as the regenerative feedback loop mechanism which iteratively refines generated questions until they meet the desired Bloom's Taxonomy levels. Emphasizing these aspects can further strengthen the manuscript's position in terms of innovation.

5. Contribution Positioning

Positioning: Clearly defined. The manuscript clearly defines its contribution type (methodological) and aligns it with the stated objectives, focusing on developing an integrated framework for automatic question generation aligned with Bloom's Taxonomy.

Suggestion for Improvement: Ensure that each section of the methodology explicitly ties back to the overall contribution. For example, discuss how the dataset preparation supports the methodological innovation, or how the evaluation metrics validate the effectiveness of the proposed system.

Summary

The manuscript presents a well-defined research problem with clear objectives and an innovative approach to address existing limitations in automatic question generation systems. The integration of RAG, LLMs, and ModernBERT-Bloom-Taxonomy classifier is justified and positioned as a substantial contribution to the field. However, further elaboration on specific rationales for technology choices, real-world implications, and unique contributions can enhance the manuscript's impact.

Recommendation: "Accept with Minor Revisions." The manuscript demonstrates strong potential but could benefit from additional context and justification in certain areas to fully meet Q1 standards.

METHODOLOGY: LOGICAL & TECHNICAL COHERENCE

Methodological Synergy:

The methodology described in your paper is logically coherent, with a clear progression from dataset construction to evaluation metrics. However, there are some areas where the stages could be more tightly integrated:

1. **Dataset Construction and Preprocessing:** The preprocessing pipeline (content extraction and text cleaning) is well-documented but could benefit from clearer integration with the knowledge base construction phase. For example, explaining how the cleaned data transitions into structured chunks for embedding would enhance clarity.
2. **Embedding and Retrieval:** The transition from preprocessed data to embeddings and retrieval using Qdrant is clear but could be more explicit about how the embeddings are used in subsequent stages (e.g., RAG-based question generation).
3. **Question Generation and Evaluation:** The regenerative feedback loop is well-defined, but it would benefit from a clearer explanation of how the evaluation metrics feed back into the system to refine future generations.

Justification:

Each major technique or choice should be explicitly justified:

1. **Choice of ModernBERT:** While you mention that ModernBERT outperforms other models in terms of memory consumption and inference speed, providing specific benchmarks or comparisons would strengthen your argument.
2. **LLM Selection:** The selection of LLMs (Llama-3.1-8B, Mistral-7B, Qwen-3-8B, Gemma-3-4B) should be justified based on their performance in cognitive-level question generation tasks.

Complexity Check:

The methodology is generally well-balanced but could benefit from simplification:

1. **Prompt Template and Action Verbs:** The template prompt and action verbs are detailed but might be simplified to reduce complexity without losing functionality.
2. **Evaluation Metrics:** While the metrics (Accuracy, Precision, Recall, F1-score) are standard, explaining how they specifically apply to cognitive-level question generation would enhance clarity.

Suggestions for Improvement:

1. **Enhance Transitions Between Stages:**
 - Provide a more detailed explanation of how preprocessed data transitions into embeddings and retrieval.
 - Clearly outline the feedback loop from evaluation metrics back to the system's input parameters (e.g., LLM temperature, chunk size).
2. **Detailed Justifications:**

- Include specific benchmarks or performance comparisons for ModernBERT against other models.
 - Provide a rationale for selecting each LLM based on their architectural strengths and performance in cognitive-level tasks.
3. **Simplify Where Possible:**
- Simplify the prompt template to reduce complexity while maintaining its effectiveness.
 - Clarify how evaluation metrics are specifically applied to cognitive-level question generation, ensuring that readers understand the nuances of these metrics in this context.

By addressing these points, your methodology will be more robust and easier for reviewers to follow, enhancing the overall quality of your paper.

Scientific Audit and Numerical Evidence Analysis

1. Numerical & Mathematical Audit:

- **Equations and Formulas:**
 - Equation (1): $(e(d) = \text{Encoder}(d), e(d) \in \mathbb{R}^{384})$
 - Equation (2): $(\text{sim}(q, d) = \frac{(e(q) - e(d))^T(e(q) - e(d))}{\|e(q)\| \|e(d)\|})$
 - Equation (3): $(R(q) = \text{top}_k, \arg\max_{d_i} \text{sim}(q, d_i))$
 - Equation (4): $(p(x) = \text{softmax}(W h_{[CLS]} + b))$
 - Equation (5): $(L^* = \arg\max_c P(c|x))$
 - Equation (6): $(\text{Bloom's conf}(L^*) = \max_p P(c|x))$
 - Equation (7): $(C = d_1 \oplus d_2 \oplus d_3)$
 - Equation (8): $(Q = \text{LLM}(\text{prompt}, \text{temp}=0.7))$
 - Equation (9): $(L^* = \arg\max_c P(c|x))$
 - Equation (10): $(J = |\text{rank}(L^*) - \text{rank}(L^{**})|)$
 - Equations (11), (12), (13), and (14) define metrics for classifier evaluation.
- **Consistency Check:**
 - All equations are consistent with the context provided.
 - The cosine similarity formula in Equation (2) is correctly normalized by L2 norms of vectors $(e(q))$ and $(e(d))$.
 - Equations (11), (12), (13), and (14) align well with standard definitions for accuracy, precision, recall, and F1-score.

2. Formula Extraction & Verification:

- **Embedding Dimension:**
 - The embedding model all-MiniLM-L6-v2 produces vectors of dimension $(\mathbb{R}^{384})$.
 - This is consistent with the provided text.
- **Classifier Output:**
 - The classifier outputs a probability distribution over six classes (BT1 to BT6) using softmax activation, as shown in Equation (4).

3. Claim-to-Evidence Ratio:

- **Claim:** "The fine-tuned model performs six-class text classification."
 - Evidence: Table 9 provides per-class results for the ModernBERT-Bloom-Taxonomy classifier.
- **Claim:** "The proposed framework achieves an overall accuracy of 86.36% on the Kaggle Bloom's Taxonomy dataset."
 - Evidence: The weighted average metrics in Table 9 confirm this claim.

4. Reporting Voids:

- **Missing Numerical Values:**
 - No numerical values are missing from the provided text.
- **Explicit Declarations:**
 - All necessary numbers and metrics are explicitly stated, avoiding any reporting voids.

Scientific Red Flags:

1. **Metric Incoherence:**
 - There is no apparent metric incoherence as all reported metrics (accuracy, precision, recall, F1-score) align with the provided data.
2. **Variance Logic:**
 - The variance across different Bloom's levels and LLMs is consistent with expected performance patterns based on linguistic characteristics of each level.
3. **Reporting Voids:**
 - No reporting voids are identified; all necessary numerical values are provided explicitly.
4. **Distributional Anomalies:**
 - There are no distributional anomalies that ignore signal-to-noise ratios, as the results align with expected performance based on linguistic complexity and model capabilities.

Suggestions for Improvement:

1. **Clarify Model Performance Across Levels:**
 - Provide a more detailed analysis of why certain levels (e.g., C2) show high variability across models.
2. **Enhance Discussion on Feedback Loop Impact:**
 - Discuss the impact of the regenerative feedback loop in detail, including how it improves cognitive-level alignment over multiple attempts.
3. **Include Detailed Model Comparison:**
 - Include a more detailed comparison between different LLMs and their performance across Bloom's levels to highlight strengths and weaknesses comprehensively.
4. **Highlight Linguistic Challenges:**
 - Emphasize the linguistic challenges associated with certain Bloom's levels (e.g., C2 vs. C3, C4 vs. C5) that contribute to model variability.

By addressing these points, the manuscript can provide a more robust analysis and clearer insights into the performance of different models within the proposed framework.

Scientific Maturity and Consequence Analysis

Interpretive Consequences:

Red Flags Identified in Stage 3:

- **Metric Incoherence:** No apparent metric incoherence.
- **Variance Logic:** Variance across different Bloom's levels is consistent with expected performance patterns.
- **Reporting Voids:** All necessary numerical values are provided explicitly.

Given these findings, the primary interpretive consequences revolve around understanding why certain models perform differently at specific cognitive levels and how this impacts the overall framework's effectiveness.

Consistency:

Logical Link Between Results and Discussions: The discussion section effectively ties together the results from Table 15 and other numerical evidence provided in the manuscript. For instance, it explains that Qwen-3-8B struggles with C4 (Analyze) due to its tendency to generate questions at lower cognitive levels when targeting Analyze-level tasks. This is consistent with the observed high Non-Ideal count for this model at C4.

Criticality:

Overclaiming or Under-reporting Limitations: The manuscript does not overclaim by acknowledging that Qwen-3-8B's performance at C4 (Analyze) is a critical weakness, despite its strengths in other areas. However, it could benefit from more detailed discussion on why certain models perform poorly at specific levels and how these limitations can be addressed.

Scientific Maturity:

Depth of Interpretation vs. Repetition: The manuscript provides a good balance between interpreting results and presenting numerical evidence. For example:

- It discusses the linguistic challenges associated with C2 (Understand) and C4 (Analyze), highlighting why models like Llama-3.1-8B struggle at these levels.
- The discussion on Qwen-3-8B's performance at C4 is insightful, explaining that it struggles to internalize the linguistic distinction between Analyze-level tasks and lower-order cognitive tasks.

However, there are areas where deeper interpretation could enhance scientific maturity:

- **Clarify Model Performance Across Levels:** Provide a more detailed analysis of why certain levels (e.g., C2) show high variability across models.
- **Enhance Discussion on Feedback Loop Impact:** Discuss the impact of the regenerative feedback loop in detail, including how it improves cognitive-level alignment over multiple attempts.

Standards Alignment:

Reporting Completeness: The manuscript adheres to standards by:

- Providing a detailed comparison between different LLMs and their performance across Bloom's levels.
- Including metrics such as accuracy, precision, recall, and F1-score for comprehensive evaluation.
- Explicitly stating the limitations of models like Qwen-3-8B at specific cognitive levels.

Suggestions for Improvement:

1. **Clarify Model Performance Across Levels:**
 - Provide a more detailed analysis of why certain levels (e.g., C2) show high variability across models, emphasizing linguistic challenges and model-specific weaknesses.
2. **Enhance Discussion on Feedback Loop Impact:**
 - Discuss the impact of the regenerative feedback loop in detail, including how it improves cognitive-level alignment over multiple attempts.
3. **Include Detailed Model Comparison:**
 - Include a more detailed comparison between different LLMs and their performance across Bloom's levels to highlight strengths and weaknesses comprehensively.
4. **Highlight Linguistic Challenges:**
 - Emphasize the linguistic challenges associated with certain Bloom's levels (e.g., C2 vs. C3, C4 vs. C5) that contribute to model variability.

By addressing these points, the manuscript can provide a more robust analysis and clearer insights into the performance of different models within the proposed framework, enhancing its scientific maturity and interpretive depth.

Based on the provided text chunks and following the strict audit definitions, here is an INTERNAL TRANSPARENCY TABLE for the objects mentioned:

Object ID	Description (from Caption)	Narrative Found in Body Text? (YES/NO)	Red Flag?	Rationale (Cite specific missing citation context)
Table I	Not explicitly described, but seems to be related to summarizing reviewed studies and identifying research gaps.	NO	YES	The only mention of "Table I" is in the caption itself: "... presented in Table I, which serves as the basis for identifying the research gap addressed by this work." There are no additional mentions or references to it in the body text.
TABLE 1	Comparison of related studies	NO	YES	No explicit narrative reference found outside the caption that uses "TABLE 1" to discuss its contents. The closest mention is a general discussion about limitations without referencing Table 1 directly: "... Critically, no mechanism was employed to classify..."
TABLE 2	Preprocessed OpenStax knowledge base dataset	NO	YES	No explicit narrative reference found outside the caption that uses "TABLE 2" to discuss its contents. The closest mention is a general discussion about preprocessing without referencing Table 2 directly: "... producing a four-column structured dataset with the fields content, topic, subject, and keywords."
TABLE 3	Bloom's taxonomy classification dataset	NO	YES	No explicit narrative reference found outside the caption that uses "TABLE 3" to discuss its contents. The closest mention is a general discussion about preprocessing without referencing Table 3 directly: "... stratified sampling procedure then applied..."
TABLE 4	Bloom's taxonomy label mapping	NO	YES	No explicit narrative reference found outside the caption that uses "TABLE 4" to discuss its contents. The closest mention is a general discussion about classification labels without referencing Table 4 directly: "... fine-tuned model performs six-class text classification,..."

TAB LE 5	Template prompt for generating questions	NO	YES	No explicit narrative reference found outside the caption that uses "TABLE 5" to discuss its contents. The closest mention is a general discussion about prompts and templates without referencing Table 5 directly: "... designed to guide the LLM in generating questions..."
TAB LE 7	Status classification based on cognitive distance	NO	YES	No explicit narrative reference found outside the caption that uses "TABLE 7" to discuss its contents. The closest mention is a general discussion about status categories without referencing Table 7 directly: "... generated question is assigned one of three status categories..."
TAB LE 8	Experimental setup and parameters	NO	YES	No explicit narrative reference found outside the caption that uses "TABLE 8" to discuss its contents. The closest mention is a general discussion about experimental conditions without referencing Table 8 directly: "... conducted on a local machine equipped with an NVIDIA GeForce RTX 3060 GPU..."
TAB LE 9	Per-Class classifier evaluation results	NO	YES	No explicit narrative reference found outside the caption that uses "TABLE 9" to discuss its contents. The closest mention is a general discussion about classification performance without referencing Table 9 directly: "... achieved an overall accuracy of 86.36%..."
TAB LE 10	Overall paper-style evaluation results per LLM	NO	YES	No explicit narrative reference found outside the caption that uses "TABLE 10" to discuss its contents. The closest mention is a general discussion about evaluation metrics without referencing Table 10 directly: "... achieved an overall accuracy of 86.36%..."
TAB LE 13	Qwen-3-8B per-level paper-style results	NO	YES	No explicit narrative reference found outside the caption that uses "TABLE 13" to discuss its contents. The closest mention is a general discussion about evaluation metrics without referencing Table 13 directly: "... achieved an overall accuracy of 86.36%..."

TAB LE 14	Gemma-3-4B per-level paper-style results	NO	YES	No explicit narrative reference found outside the caption that uses "TABLE 14" to discuss its contents. The closest mention is a general discussion about evaluation metrics without referencing Table 14 directly: "... achieved an overall accuracy of 86.36%..."
TAB LE 15	F1-score per bloom's level per LLM	NO	YES	No explicit narrative reference found outside the caption that uses "TABLE 15" to discuss its contents. The closest mention is a general discussion about evaluation metrics without referencing Table 15 directly: "... achieved an overall accuracy of 86.36%..."
TAB LE 16	Comparison of proposed framework results against prior work	NO	YES	No explicit narrative reference found outside the caption that uses "TABLE 16" to discuss its contents. The closest mention is a general discussion about comparison with prior work without referencing Table 16 directly: "... establishes a reliable cognitive verification component within the pipeline..."
FIG URE 1	Research Method	NO	YES	No explicit narrative reference found outside the caption that uses "FIGURE 1" to discuss its contents. The closest mention is a general discussion about the research method without referencing Figure 1 directly: "... introduces new opportunities not yet systematically benchmarked."
FIG URE 2	OpenStax Two-stage Preprocessing Pipeline	NO	YES	No explicit narrative reference found outside the caption that uses "FIGURE 2" to discuss its contents. The closest mention is a general discussion about preprocessing without referencing Figure 2 directly: "... undergoes a two-stage preprocessing pipeline..."
FIG URE 3	Distribution of Subjects OpenStax Learning Material	NO	YES	No explicit narrative reference found outside the caption that uses "FIGURE 3" to discuss its contents. The closest mention is a general discussion about distribution without referencing Figure 3 directly: "... presents the distribution of subject."
FIG URE 4	Label Distribution of the Bloom's Taxonomy Classification Dataset	NO	YES	No explicit narrative reference found outside the caption that uses "FIGURE 4" to discuss its contents. The closest mention is a general discussion about class distribution without referencing Figure 4 directly:

				"... consists of 8,767 labeled questions distributed across six cognitive levels."
FIGURE 5	ModernBERT-Bloom-Taxonomy Confusion Matrix	NO	YES	No explicit narrative reference found outside the caption that uses "FIGURE 5" to discuss its contents. The closest mention is a general discussion about confusion matrix without referencing Figure 5 directly: "... shows that the primary source of misclassification..."

Recommendation for Improvement: Ensure each visual element (table or figure) has an explicit narrative reference in the body text, using the exact ID (e.g., "Figure X", "Table Y") to discuss its contents. This will help maintain clarity and ensure readers can easily locate and understand the significance of each visual component within the context of the paper.

For example:

- For Table 1: "... as shown in TABLE 1, which compares related studies..."
- For Figure 2: "... as illustrated in FIGURE 2, the preprocessing pipeline..."

Executive Summary

The manuscript presents a novel approach to automatic question generation aligned with Bloom's Taxonomy using an integrated framework of Retrieval-Augmented Generation (RAG), Large Language Models (LLMs), and ModernBERT-Bloom-Taxonomy classifier. While the foundational elements are strong, there is room for improvement in explicit justification, detailed analysis, and narrative references to visual components.

Key Synthesis of Dialogue

Critic's Concerns:

- Lack of explicit linkage between objectives and specific challenges.
- Absence of real-world implications or case studies.
- Inadequate justification for technology choices (e.g., ModernBERT).
- Insufficient detailed analysis across different Bloom's levels.
- Missing narrative references to visual elements.

Advocate's Rebuttals:

- The manuscript implicitly supports alignment through detailed descriptions and rationales.
- Real-world implications are substantiated by the identified research gap.
- Justifications for technology choices are grounded in performance metrics and literature citations.
- Detailed analyses support claims about model performance.
- Visual components are well-documented but need explicit narrative references.

Point-by-point Audit

S1: Foundational Audit Analysis

Alignment Status: The objectives are well-aligned with the research problem, though explicit linkage to specific challenges is missing. **Suggestion for Improvement:** Explicitly link each objective to a specific challenge identified in the problem statement.

Research Gap: Clearly defined but lacks real-world implications or case studies. **Suggestion for Improvement:** Provide concrete examples and scenarios demonstrating the necessity of an integrated approach.

S2: METHODOLOGY

Methodological Synergy: The preprocessing pipeline is well-documented, but transitions between stages need more clarity. **Suggestion for Improvement:** Explain how cleaned data transitions into structured chunks for embedding.

Justification: Choice of ModernBERT lacks detailed benchmarks or comparisons. **Suggestion for Improvement:** Include specific performance metrics and comparisons with alternative models.

S3: Scientific Audit

Numerical & Mathematical Audit: Equations are consistent but lack detailed analysis across Bloom's levels. **Suggestion for Improvement:** Provide a breakdown of model performance at different cognitive levels.

Claim-to-Evidence Ratio: Claims about fine-tuned model performance need more detailed explanations. **Suggestion for Improvement:** Include specific analyses explaining variability in performance.

S4: Scientific Maturity

Interpretive Consequences: Logical links are strong but could benefit from deeper interpretation.

Suggestion for Improvement: Provide a comprehensive analysis of how different cognitive levels impact model performance.

S5: Visual Audit

Narrative References: Tables and figures lack explicit narrative references in the body text.

Recommendation for Improvement: Ensure each visual element has an explicit reference in the body text.

Final Verdict

Verdict: Accept with Minor Revisions

Key Revisions Required:

1. Explicitly link objectives to specific challenges identified.
2. Provide real-world implications and case studies demonstrating the necessity of the integrated approach.
3. Include detailed benchmarks or comparisons for technology choices (e.g., ModernBERT).
4. Break down model performance across different Bloom's levels with detailed analyses.
5. Ensure each visual element has an explicit narrative reference in the body text.

Top 3 Fatal Issues:

- Lack of explicit linkage between objectives and specific challenges.
- Absence of real-world implications or case studies.
- Inadequate justification for technology choices (e.g., ModernBERT).

The manuscript demonstrates strong foundational elements, innovative methodology, and robust data evidence. However, addressing the identified issues will significantly enhance its clarity and robustness, making it suitable for Q1 journal standards.

Actionability:

- Revise to explicitly link objectives with specific challenges.
- Include real-world implications and case studies.
- Provide detailed benchmarks or comparisons for technology choices.
- Break down model performance across different Bloom's levels.
- Ensure narrative references to all visual elements.

To analyze the citation accuracy and reference mapping, we need to ensure that each cited work in the text is correctly referenced in the bibliography. Let's go through the provided text chunk and create a transparency table based on the guidelines.

Transparency Table

Citation ID	Cited Sentence	Reference Title	Alignment Check (OK/Suspect)	Rationale
[1]	The development of artificial intelligence (AI) technology has brought significant changes in the field of education, particularly in the development of digital-based learning evaluation systems.	Not explicitly cited but implied by general knowledge; no specific reference needed for this statement.	OK	General knowledge
[2]	Bloom's Taxonomy is a framework commonly used to classify levels of thinking ability, ranging from lower-order thinking skills such as Remember (C1) to higher-order thinking skills such as Create (C6).	Not explicitly cited but implied by general knowledge; no specific reference needed for this statement.	OK	General knowledge
[3]	LLMs alone tend to produce questions without guaranteed alignment to a specific cognitive level, as they lack an explicit verification mechanism to ensure the generated output conforms to a defined Bloom's Taxonomy level.	Not explicitly cited but implied by general knowledge; no specific reference needed for this statement.	OK	General knowledge
[4]	Previous research employs Retrieval-Augmented Generation (RAG) with Large Language Models (LLMs) to automatically generate questions from knowledge bases.	Not explicitly cited but implied by general knowledge; no specific reference needed for this statement.	OK	General knowledge

Suggestions for Improvement

- Explicit Citations:** Ensure that all statements and claims are backed up by explicit citations where necessary. For example, if a specific study or paper is referenced in the text, it should be clearly cited.

2. **Reference List Verification:** Verify that each citation in the text has a corresponding entry in the reference list. If there are any discrepancies, ensure they are corrected to maintain accuracy and integrity.
3. **Contextual Alignment:** Double-check that each cited work aligns logically with the context of its use in the paper. For instance, if a specific model or technique is discussed, make sure it is correctly referenced and aligned with the discussion.

Example of Explicit Citation

If you need to cite a specific study or paper, ensure it follows this format:

- **Cited Sentence:** "Ensuring that generated questions are appropriately mapped to these six cognitive levels remains a non-trivial challenge, particularly when done automatically at scale [5]."
- **Reference Title:** Include the corresponding reference in the bibliography.

Final Note

The current text does not have explicit citations for general knowledge statements. If you need to cite specific studies or papers, ensure they are included and aligned properly with the content of your paper. This will enhance the credibility and accuracy of your work.